Principled Private Data Release with Deep Learning

Michael Fine Harvard University
Cambridge MA, 02138
`mfine@college.harvard.edu`

November 18, 2019

1 Introduction

2 Related Work

- [JY19]
- [GLL$^+$17]
- [NRVW19]
- [AGH18]
- [NRW18]
- [GAH$^+$14]
- [HLM12]

3 Background

3.1 Query Release Problem

We study the problem of privately generating synthetic data to answer statistical queries over a data universe $\mathcal{X}$. Formally, a statistical query over $\mathcal{X}$ is a function $q : \mathcal{X} \to \{0,1\}$. Given a a dataset $x \in \mathcal{X}^n$, we define $q(x) = \sum_{x_i \in x} q(x_i)$.

Our goal is to produce a synthetic dataset that, for every query in some family of queries, takes approximately the same value as the true dataset.

Definition 1 ($\alpha-$approximate)**.** *We say a synthetic dataset x $\alpha-$approximates a true dataset $\hat{x}$ w.r.t a family of statistical queries $\mathcal{Q}$ if*

$$\forall q \in \mathcal{Q} : \quad |q(x) - q(\hat{x})| \leqslant \alpha \tag{1}$$

[TODO finish up]

3.2 Game Theoretic Formulation

One can formulate the problem of producing an $\alpha-$approximate dataset as a two-player, zero sum game [HRU13] between a discriminator D and a generator G. The generator has an action set $\mathcal{X}$, while the discriminator has an action set $\mathcal{Q}$. The generator aims to output a dataset $x \in \mathcal{X}$ that maximally agrees with $\hat{x}$, while the discriminator aims to find queries $q \in \mathcal{Q}$ that distinguish $\hat{x}$ and x.

Formally, given a play $x \in \mathcal{X}$ and $q \in Q$, the discriminator gets payoff $V(x, q)$ and the generator gets payoff $-V(x, q)$, where $V(x, q)$ denotes:

Definition 2 (Payoff).
$$V(x, q) := |q(x) - q(\hat{x})| \tag{2}$$

The goal of both G and D is to maximize their worst case payoffs, thus

$$\max_{q \in \mathcal{Q}} \min_{x \in X} V(x, q) \quad \text{(Goal of } D\text{)} \quad and \quad \min_{x \in X} \max_{q \in \mathcal{Q}} V(x, q) \quad \text{(Goal of } G\text{)} \tag{3}$$

If there exists a point (x^*, q^*) such that neither G nor D can improve their payoffs by playing a different move, we call that a *Pure Nash Equilibrium*. Unfortunately, a pure equilibrium is not always guaranteed to exist (and likely does not in the case of synthetic data generation).

However, the seminal work of Nash et. al showed that there always exists a *Mixed Nash Equilibrium (MNE)*, where the players play *probability distributions* over their action sets, instead of fixed actions.

Let $\Delta(\mathcal{X})$ and $\Delta(\mathcal{Q})$ denote the set of probability distribution over $\mathcal{X}$ and $\mathcal{Q}$. Formally, if G plays a strategy $g \in \Delta(\mathcal{X})$ and D plays $d \in \Delta(\mathcal{Q})$, we define the payoff to be the expected value of a single draw:

$$V(g, d) := \mathbb{E}_{x \sim g, q \sim d} V(x, q) \tag{4}$$

Thus, a pair of strategies $g \in \Delta(\mathcal{X})$ and $d \in \Delta(\mathcal{Q})$ forms an $\alpha-$approximate mixed nash equilibrium if for all strategies $u \in \Delta(\mathcal{X})$ and $w \in \Delta(\mathcal{Q})$

$$V(g, w) \leq V(u, w) + \alpha \quad and \quad V(u, d) \leq V(u, w) - \alpha \tag{5}$$

Moreover, Gaboardi et. al showed how to reduce the problem of finding an $\alpha-$approximate dataset to the problem of finding an an $\alpha-$equilibrium in the query release game:

Theorem 1. *Let (u, w) be the $\alpha-$approximate MNE in a query release game for a dataset $\hat{x} \in \mathcal{X}$ and a query universe $\mathcal{Q}$. If $\mathcal{Q}$ is closed under negation, then the dataset S sampled from u α-approximates $\hat{x}$ over Q. [GAH$^+$14]*

Hence, our task is to provide an algorithm to private reach an $\alpha-$MNE in the query release game. In the following section, we will provide the background for how this can be done with GANs.

3.3 Online Learning

To efficiently find equilibrium in the zero-sum GAN game, we draw on results from online learning. In the online learning setting, in each of T rounds a player is given a loss function f_t , possibly adversarial chosen. The players goal is to chose an action $x_{t+1} \in \mathcal{X}$ in order to minimize the cumulative *regret*.

Definition 3 (Regret). *The regret measures the cumulative loss of the player, compared to the best fixed decision in hindsight.*

$$Regret_T(f_1, ..., f_T) = \sum_{t=1}^{T} f_t(x_t) - min_{x \in \mathcal{X}} \sum_{t=1}^{T} f_t(x^*) \tag{6}$$

When a strategy provably leads to regret is sublinear in T, we call that *no-regret*, as regret $\to 0$ as $T \to \infty$. One approach to regret minimization is to choose the action x_{t+1} that minimizes the cumulative loss over all past loss functions l

$$x_{t+1} = \arg\min_{x \in \mathcal{X}} \sum_{t=1}^{T} f_t(x) \tag{7}$$

This approach is known as *Follow-The-Leader*. While natural, this approach is easily exploitable by an adversary. At a high level, this is because it *overfits* to past outcomes, allowing it to optimize between suboptimal strategies. To rectify this, a powerful strategy is *Follow-the-Regularized-Leader*

Definition 4 (Follow-The-Reguralized-Leader (FTRL)). *Given a reguralization function $R(x)$ and a regularization weight η_T, at each step choose $x_{t+1} \in \mathcal{X}$ to minimize the regularized cumulative loss:*

$$x_{t+1} = \arg\min_{x \in \mathcal{X}} \sum_{t=1}^{T} +\eta_T R(X) \tag{8}$$

One common regularization function is the l_2 norm $R(x) = ||x||_2$.

When the loss function f_T is convex, FTRL can be shown to be no-regret [Haz19]. Unfortunately, in the GAN setting, where the loss function f_T is defined by a highly non-convex deep network, we don't have that guarantee.

3.3.1 Online Learning for Non-convex losses

In general, finding the minimum of a sum of non-convex functions is hard. However, in practice gradient descent over neural networks has proven to be remarkably effective at approximately solving non-convex loss functions. This leads to the natural question: assuming we have an offline non-convex optimization oracle $\mathcal{O}$, can we use that to find a no-regret strategy in the online setting? Agarwal et al showed that we can, using a Follow-The-Leader variant known as Follow-The-Perturbed-Leader [AGH18]. Formally:

Definition 5 (Offline optimization oracle). *Let $\mathcal{O}$ take a sequence of (possibly non-convex) loss functions $(f_1...f_T) \in \mathcal{L}^T$ and a d-dimensional vector d, and output $x^* \in \mathcal{X}$*

$$x^* \in \arg\min_{x \in \mathcal{X}} \sum_{t=1}^{T} f_t(x) + \sigma^T x \tag{9}$$

If we relax the requirement to allow $\mathcal{O}$ to output an approximate minimizer x^*

$$x^* \leq \arg\min_{x \in \mathcal{X}} \left(\sum_{t=1}^{T} f_t(x) + \sigma^T x \right) + \alpha \tag{10}$$

we then call $\mathcal{O}$ an $\alpha-$approximate offline oracle.

We can use this offline oracle $\mathcal{O}$ to minimize regret in the online case:

Definition 6 (Follow-The-Perturbed-Leader (FTPL)). *Given an offline oracle $\mathcal{O}$ and a parameter η, at each step t FTPL draws a random vector $\sigma_t \sim Exp(\eta)^d$. It then outputs*

$$x^* \in \arg\min_{x \in \mathcal{X}} \sum_{t=1}^{T} f_t(x) + \sigma^T x \tag{11}$$

Theorem 2. *FTPL has sublinear regret. [AGH18]*

Building on the work of Freund and Schapire, we can show that if G and D both play FTPL for T rounds, they will converge to an $\alpha-$approximate equilibrium. Formally:

Theorem 3. *Suppose that G and D play according to FTPL. We can choose $T \in poly(d)/\alpha^3$ such that the expected average regret of FTPL is at most α. Then, G_T and D_T produce strategies in an $\alpha-$approximate equilibrium [AGH18].*

3.4 Generative Adversarial Networks

Generative Adversarial Networks (GANs), introduced by Goodfellow et. al, are an approach to generative deep learning that has shown remarkable promise in generating high fidelity samples [GPM+14]. In the GAN setup, a generator G is paired with a discriminator D. At each round, D is trained to distinguish real samples drawn from P_{data} from generated samples drawn from P_g, while G is trained to generate realistic samples that fool the discriminator.

This yields a two player, zero sum game with minimax objective

$$\min_G \max_D V_{gan}(G, D) := \frac{1}{2}\mathbb{E}_{\mathbf{x} \sim p_{data}} \log D(\mathbf{x}) + \frac{1}{2}\mathbb{E}_{z \sim p_z} \log(1 - D(G(z))) \tag{12}$$

However, [ACB17] showed that this cost function V_{gan} is not sensible cost function in practice, when the distributions are supported by low-dimensional manifolds. Instead, they proposed to use the Earth-Mover, or Wasserstein-1 objective.

Definition 7 (Earth Mover Distance). *The EM distance between two distribution $\mathbb{P}_r$ and $\mathbb{P}_g$ is*

$$W(\mathbb{P}_r, \mathbb{P}_g) := \inf_{\gamma \in \Pi(\mathbb{P}_r, \mathbb{P}_g)} \mathbb{E}_{(x,y) \sim \gamma} |x - y| \tag{13}$$

where $\Pi(\mathbb{P}_r, \mathbb{P}_g)$ denotes the set of all joint distributions $\gamma(x, y)$ whose marginals are respectively $\mathbb{P}_r$ and $\mathbb{P}_g$.

While this infimum is highly intractable to compute, [ACB17] used the Kantorovich-Rubinstein duality [Vil08] to show that

$$W(\mathbb{P}_r, \mathbb{P}_g) = \sup_{\|f\|_L \leqslant 1} \mathbb{E}_{x \sim \mathbb{P}_r}[f(x)] - \mathbb{E}_{x \sim \mathbb{P}_g}[f(x)] \tag{14}$$

where the supremum is over all 1-Lipschitz functions.

Remark 1. *Note the remarkable similarity of the GAN objective V_{gan} to the earlier query release objective V. Recall that the mixed strategy V is defined as*

$$\min_{g \in \mathcal{G}} \max_{d \in \mathcal{D}} V(g,d)^2 := \mathbb{E}_{x \sim g, q \sim d} V(x,q)$$
$$= \mathbb{E}_{x \sim g, q \sim d} |q(x) - q(\hat{x})|$$
$$= \mathbb{E}_{x \sim g, \hat{x} \sim p_{data}} |D(x) - D(\hat{x})|$$
$$= \mathbb{E}_{x \sim g, \hat{x} \sim p_{data}} [D(x) - D(\hat{x})|D(x) \geqslant D(\hat{x})]$$

Further, we assume that

4 QueryGAN

Theorem 3 requires the existence of an actual oracle $\mathcal{O}$ that can find the minimum of a perturbed sum of non-convex functions. In this case, we need an oracle that can minimize the sub of deep neural networks. While SGD is remarkably effective in practice, we are unable to provide guarantees (probabilistic or otherwise) about how close the convergent solution SGD outputs is to the global minima. Recent work suggests that this is not a problem in practice, as spurious local minima (local minima significantly worse than the global minima) get exponentially rarer as the network gets larger [CHM+14]. However, these results still rely on too many impractical assumptions to make them relevant in practice.

However, while we cannot guarantee (or even certify) convergence to an approximate global minima in general, we can take advantage of the specific structure of the query release problem.

<div style="border:1px solid orange">summarize approach</div>

<div style="border:1px solid orange">Show that GAN objective is concave wrt discriminator paramters ala [GLL+17]</div>

Theorem 4. *For any discriminator, the optimal generator G* has payoff $-V(g,d) = 0$. One such generator draws a uniform sample from the true dataset $\hat{x}$.*

Proof. Recall that the generators payoff is $-V(g,d)$, where the value V is defined as $V(g,d) := \mathbb{E}_{x \sim g, q \sim d} |q(x) - q(\hat{x})|$, where $\hat{x}$ is the true, sensitive dataset. Clearly, the generator's payoff is at most 0. This is trivially attainable if g is the uniform distribution over the rows of $\hat{x}$. □

Theorem 4 allows us to track how close to optimal the generator G_i is at each step. Assuming D_i is optimal, the regret at each step is simply the generator loss $V(G_i, D_i)$. The cumulative regret is simply the sum of the generator loss at each step.

This gives us a way to track *generator* optimality, assuming the discriminator is optimal. Unlike the generator, there is no obvious maximum payoff for the discriminator in general. However, if we restrict the class of discriminators to single layer neural networks parametrized by θ:

$$D_\theta(x) = \sigma(\theta^T x + b) \tag{15}$$

then $V(G, D_\theta)$ becomes convex with respect to θ, and we can use standard gradient descent methods to guarantee convergence to a global optimum.

At each round, we update D and G accoring to gradient descent over FTPL. For convenience, let the notation $f_{0:t}$ denote the set of all functions $(f_0, ..., f_t)$.

<div style="border:1px solid orange">Make clearer the distinction between queries over single rows and queries over all rows,</div>

<div style="border:1px solid orange">Maybe another line proving this</div>

<div style="border:1px solid orange">Prove gradient descent bounds</div>

$$\mathcal{O}_d(\theta_t, f_{0:t}) := \theta_t - \nabla_{\theta_t} \sum_{t=1}^{T} f_t + \sigma_1^T x \tag{16}$$

$$\mathcal{O}_g(\phi_t, g_{0:t}) := \phi_t - \nabla_{\theta_t} \sum_{t=1}^{T} g_t + \sigma_2^T x \tag{17}$$

Algorithm 1: QueryGAN

Input: one-layer discriminator D_θ, deep generator G_ϕ, discriminator and generator oracle $\mathcal{O}_d, \mathcal{O}_g$, Rounds T, noise η, output dimension d, game objective V

Result: Trained generator G_T, Accuracy α

for $t \in 1...T$ **do**

 Draw discriminator and generator perturbations
 $\sigma_1 \sim Exp(\eta)^d \quad and \quad \sigma_2 \sim Exp(\eta)^d$

 Update D and G with their respective oracles:
 $\theta_{t+1} \leftarrow \mathcal{O}_d(\theta_t, f_{0:t}) \quad and \quad \phi_{t+1} \leftarrow \mathcal{O}_d(\phi_t, g_{0:t})$

 Update losses:
 $f_{t+1}(\cdot) = V(\cdot, D_{\phi_{t+1}}) \quad and \quad g_{t+1}(\cdot) = V(G_{\theta_{t+1}}, \cdot)$

Calculate cumulative regret: $R \leftarrow \sum_{t \in T} f_t(G_{\phi_t})$
return *Mixed strategy:* $G \sim Unif\{G_{\theta_1}, ..., G_{\theta_T}\}$, *Regret:* R

Theorem 5. *Let G, α be the results of running QueryGAN with inputs . Let x be the dataset sampled from G*

$$x := \{G(z_1), ..., G(z_n)\} \quad where \quad z \sim \mathcal{N}_n(1, I_n) \tag{18}$$

 Then x is $\alpha-$approximate with respect to all queries representable by D.

Proof. □

4.1 QueryGAN privacy

Privacy is ensured by the addition of exponentially distributed noise σ_1, σ_2. Interestingly, the original purpose of this noise is not privacy, but rather to ensure convergence of the online algorithm. Because of the deep connections between differential privacy and online learning [NRVW19] [GHM19], however, this noise also ensures ϵ-differential privacy.

4.1.1 Exponential noise

The standard mechanism for ensuring differential privacy is the Laplace mechanism:

Definition 8 (Laplace Mechanism [DR13])**.** *Given any function $f : \mathbb{N}^{|\mathcal{X}|} \to \mathbb{R}^k$, the laplace mechanism is defined as*

$$\mathcal{M}_L(x, f(\cdot), \epsilon) = f(x) + (Y_1, \ldots, Y_k) \tag{19}$$

 where each Y_i are i.i.d drawn from $Lap(\nabla f / \epsilon)$

It's easy to show that the Laplace mechanism preserves $(\epsilon, 0)$-differential privacy. However, in *QueryGAN* we add noise drawn from the exponential distribution:

Definition 9 (Exponential Distribution). *The exponential distribution with parameter λ is the distribution with density function*

$$Exp(x; \lambda) := \begin{cases} \lambda e^{-\lambda x} & x \geqslant 0 \\ 0 & x < 0 \end{cases} \tag{20}$$

Note that the Laplace mechanism can be seen as the symmetric form of the exponential distribution. Specifically, if $X \sim Lap(\lambda)$, then $|X| \sim Exp(1/\lambda)$.

4.1.2 Tracking privacy loss with moments accountant

use [NRVW19] to show privacy

4.1.3 Reporting Regret Bounds

GAN privacy

-

-

-

Talk about privately reporting α with report noisy max

4.2 Extensions

4.2.1 Marginals

While constraining D to be a one layer linear discriminator is restricting, it still is capable of representing a number of query families of interest. Specifically, D contains all k-way marginals.

Maybe PATE-GAN

Definition 10 (Marginal). *A marginal $m : \mathcal{X} \rightarrow \{0, 1\}$ over a row $x \in \{0,1\}^m$ is a monotone conjunction, parametrized by some subset S of the input features.*

$$m_S(x) = \prod_{i \in S} x_i \tag{21}$$

We extend this to a dataset X of n rows by defining $m(X) = \sum_{x \in X} m(x)$. A k-way marginal is a marginal restricted to k features.

[DR13]

A k-way marginal can be thought of as counting the number of rows with the same value in the chosen k features. Marginals are a useful way of providing a synopsis of a dataset that still captures complex relationships between features. Producing a differentially private synthetic dataset that agrees with all $k-$way marginals of the true dataset is an extremely well studied problem in the field . .

However, we can show that if *QueryGAN* succeeds, it is able to match all $k - way$ marginals. This follows from the fact that a linear discriminator can contain all marginals.

Survey marginal results

This theorem relies on the use of a Rectified Linear Unit activation function

Impossibility results

Definition 11 (ReLU). *The ReLU activation function $R(x) : \mathbb{R} \rightarrow \mathbb{R}^+ := max(0, x)$*

This non-linearity allows us to approximate the nonlinear marginal query with a linear neural network:

Theorem 6. *Let D be single layer discriminator parametrized by θ with a ReLU activation function s.t. $D_\theta(x) = \sigma(\theta^T x + b)$. For any single-row marginal m, there exists θ, b s.t. $D_\theta(x) = m(x)$ for all x.*

Proof. This follows from the definition of a marginal. Let m_S be the marginal over the features S. Let $\theta_i = \mathbb{1}_{i \in S}$. Setting $b = 1 - |S|$, it's clear that

$$D_\theta(x) + 1 = \begin{cases} 1, & \prod_{i \in S} x_i = 1 \\ 0, & \text{otherwise} \end{cases} = m_S(x) \tag{22}$$

$\square$

Theorem 7. *Let G, α be the output of running QueryGAN with __. Then G will generate a__ dataset with all marginal counts accurate to within α.*

Parameters, steps etc

The proof of this theorem follows directly from Theorem 5 and Theorem 6.

Oracle runtime (also comment on non-oracle runtime)

4.3 QueryGAN with alternate discriminators

While QueryGAN follows the standard GAN practice of representing both the generator and discriminator with a neural network, this is not mandatory. Indeed, given that D is represented by the almost trivially simple 1 layer network, D is best understood in more general terms than a neural network.

Summary, contextualize empirical results

Definition 12 (Tractable Discriminator Set). *We say a class of functions $\mathcal{F} : \mathcal{X} \to \{0, 1\}$ parametrizes a set of tractable discriminators w.r.t a class of queries $\mathcal{Q}$ iff*

1. *$\mathcal{Q} \subseteq \mathcal{F}$*

2. *There exists a tractable offline oracle $\mathcal{O}$* _____

What does it do (also define tractable)

As shown above, the set $\mathcal{F}_{single}$ of one layer neural networks is a tractable discriminator set for all marginals (as well as all sigmoided linear functions in general). This has the benefit of being easy to optimize in practice, without the runtime depending exponentially on the dimensionality of the query space. Relaxing that restriction by allowing for less efficiently optimizable discriminators lets us generate α-accurate synthetic data for much larger query classes.

4.3.1 Multiplicative Weights

Consider the application of the renowned Multiplicative Weights algorithm to the query release problem, introduced in [HR10].

Definition 13 (Multiplicative Weight Oracle Algorithm). *Fix a class of queries $\mathcal{Q}$ and a true dataset $\hat{x}$. Define an initial uniform distribution over queries θ_0. Let $\mathcal{O}_{MW}(\theta_t, G_t)$ output a reweighted distribution*

$$\theta_{t+1}^q \propto \exp\left(-\eta V(G_t, q)\right) \cdot \theta_t^q \tag{23}$$

for each $q \in \mathcal{Q}$

θ_t defines a distribution, with each query weighted proportionally to how well it distinguishes real data from fake data. The discriminator draws $D_{\theta_t}(x) := |q(x) - q(\hat{x})|$ where q is a query drawn $q \sim \theta_t$.

4.4 TODO

- Local DP GAN

- FTPL oracle comes with privacy for free – but also talk about how to handle non-private oracle

Proofs:

- GAN objective =? query release objective

- FTPL gradient descent (not clear this is necessarily a full on proof)

- Proof of accuracy of subsampled FTPL objective

4.5 Adversarial Kernel Learning

Limiting

5 Empirical Results

6 Conclusion

References

[ACB17]   Martin Arjovsky, Soumith Chintala, and Léon Bottou. Wasserstein GAN. January 2017.

[AGH18]   Naman Agarwal, Alon Gonen, and Elad Hazan. Learning in Non-convex Games with an Optimization Oracle. *arXiv:1810.07362 [cs, stat]*, October 2018.

[CHM$^+$14] Anna Choromanska, Mikael Henaff, Michael Mathieu, Gérard Ben Arous, and Yann LeCun. The Loss Surfaces of Multilayer Networks. *arXiv:1412.0233 [cs]*, November 2014.

[DR13]   Cynthia Dwork and Aaron Roth. The Algorithmic Foundations of Differential Privacy. *Foundations and Trends® in Theoretical Computer Science*, 9(3-4):211–407, 2013.

[GAH$^+$14] Marco Gaboardi, Emilio Jesús Gallego Arias, Justin Hsu, Aaron Roth, and Zhiwei Steven Wu. Dual Query: Practical Private Query Release for High Dimensional Data. *arXiv:1402.1526 [cs]*, February 2014.

[GHM19]   Alon Gonen, Elad Hazan, and Shay Moran. Private Learning Implies Online Learning: An Efficient Reduction. *arXiv:1905.11311 [cs, stat]*, May 2019.

[GLL$^+$17] Paulina Grnarova, Kfir Y. Levy, Aurelien Lucchi, Thomas Hofmann, and Andreas Krause. An Online Learning Approach to Generative Adversarial Networks. *arXiv:1706.03269 [cs, stat]*, June 2017.

[GPM+14] Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative Adversarial Nets. In Z. Ghahramani, M. Welling, C. Cortes, N. D. Lawrence, and K. Q. Weinberger, editors, *Advances in Neural Information Processing Systems 27*, pages 2672–2680. Curran Associates, Inc., 2014.

[Haz19] Elad Hazan. Introduction to Online Convex Optimization. *arXiv:1909.05207 [cs, math, stat]*, September 2019.

[HLM12] Moritz Hardt, Katrina Ligett, and Frank Mcsherry. A Simple and Practical Algorithm for Differentially Private Data Release. In F. Pereira, C. J. C. Burges, L. Bottou, and K. Q. Weinberger, editors, *Advances in Neural Information Processing Systems 25*, pages 2339–2347. Curran Associates, Inc., 2012.

[HR10] M. Hardt and G. N. Rothblum. A Multiplicative Weights Mechanism for Privacy-Preserving Data Analysis. In *2010 IEEE 51st Annual Symposium on Foundations of Computer Science*, pages 61–70, October 2010.

[HRU13] Justin Hsu, Aaron Roth, and Jonathan Ullman. Differential Privacy for the Analyst via Private Equilibrium Computation. *Proceedings of the 45th annual ACM symposium on Symposium on theory of computing - STOC '13*, page 341, 2013.

[JY19] James Jordon and Jinsung Yoon. PATE-GAN: GENERATING SYNTHETIC DATA WITH DIFFERENTIAL PRIVACY GUARANTEES. page 21, 2019.

[NRVW19] Seth Neel, Aaron Roth, Giuseppe Vietri, and Zhiwei Steven Wu. Differentially Private Objective Perturbation: Beyond Smoothness and Convexity. *arXiv:1909.01783 [cs, stat]*, September 2019.

[NRW18] Seth Neel, Aaron Roth, and Zhiwei Steven Wu. How to Use Heuristics for Differential Privacy. *arXiv:1811.07765 [cs, stat]*, November 2018.

[Vil08] Cédric Villani. *Optimal Transport: Old and New*, volume 338. Springer Science & Business Media, 2008.